

```
1 %load_ext autoreload
2 %autoreload 2
3
4 from google.colab import drive
5 drive.mount('/content/drive')
```

➞ Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_

Enter your authorization code:

.....

Mounted at /content/drive



```
1 import os, sys
2
3 DATAPATH = '/content/drive/My Drive/Coursera/EDHEC/investment-portfolio/data'
4 print(f"DATAPATH:{DATAPATH} contents:{os.listdir(DATAPATH)}")
5
6 MODULEPATH = '/content/drive/My Drive/Coursera/EDHEC/investment-portfolio/nb'
7 print(f"MODULEPATH:{MODULEPATH} contents:{os.listdir(MODULEPATH)}")
8
9 sys.path.append(MODULEPATH)
10 print(f"sys.path:{sys.path}")
```

➞ DATAPATH:/content/drive/My Drive/Coursera/EDHEC/investment-portfolio/data contents:
MODULEPATH:/content/drive/My Drive/Coursera/EDHEC/investment-portfolio/nb contents:
sys.path:['', '/env/python', '/usr/lib/python36.zip', '/usr/lib/python3.6', '/usr/lib/python3.6/lib-dynload', '/usr/local/lib/python3.6/dist-packages', '/usr/lib/python3.6/dist-packages']



▼ Deviations from Normality

(plus python functions with default parameters plus a quick example of recursive functions)

Today, we'll develop the code for skewness and kurtosis even though these are already available in we'll apply them to hedge fund index returns.

We'll also look at using `scipy.stats` module to apply the *Jarque-Bera* test for normality, and apply

First, add the following code to our `edhec_risk_kit.py`

```
def get_hfi_returns():
    """
    Load and format the EDHEC Hedge Fund Index Returns
    """
    hfi = pd.read_csv("data/edhec-hedgefundindices.csv",
                      header=0, index_col=0, parse_dates=True)
    hfi = hfi/100
    hfi.index = hfi.index.to_period('M')
    return hfi
```

```

1 %load_ext autoreload
2 %autoreload 2
3
4 import pandas as pd
5 import edhec_risk_kit_105_BBI as erk
6 hfi = erk.get_hfi_returns(DATAPATH)
7 hfi.head()

```

➞ The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

| | Convertible Arbitrage | CTA Global | Distressed Securities | Emerging Markets | Equity Market Neutral | Event Driven | Fixed Income Arbitrage | Glo Ma |
|---------|--------------------------|---------------|--------------------------|---------------------|-----------------------------|-----------------|------------------------------|-----------|
| date | | | | | | | | |
| 1997-01 | 0.0119 | 0.0393 | 0.0178 | 0.0791 | 0.0189 | 0.0213 | 0.0191 | 0.0 |
| 1997-02 | 0.0123 | 0.0298 | 0.0122 | 0.0525 | 0.0101 | 0.0084 | 0.0122 | 0.0 |
| 1997-03 | 0.0078 | -0.0021 | -0.0012 | -0.0120 | 0.0016 | -0.0023 | 0.0109 | -0.0 |
| 1997-04 | 0.0086 | -0.0170 | 0.0030 | 0.0119 | 0.0119 | -0.0005 | 0.0130 | 0.0 |
| 1997-05 | 0.0156 | -0.0015 | 0.0233 | 0.0315 | 0.0189 | 0.0346 | 0.0118 | 0.0 |



▼ Skewness

Intuitively, a negative skew means that you get more negative returns than you would have expected from the normal distribution.

Another way of thinking about it is if that returns are normally distributed, the mean and the median are equal.

However, if they are negatively skewed, the expected value i.e. the mean is less than the median. If they are positively skewed, the expected value (again, the mean) is greater than the median.

```
1 pd.concat([hfi.mean(), hfi.median(), hfi.mean(>hfi.median())], axis=1)
```



| | 0 | 1 | 2 |
|-------------------------------|-----------|---------|-------|
| Convertible Arbitrage | 0.005508 | 0.0065 | False |
| CTA Global | 0.004074 | 0.0014 | True |
| Distressed Securities | 0.006946 | 0.0089 | False |
| Emerging Markets | 0.006253 | 0.0096 | False |
| Equity Market Neutral | 0.004498 | 0.0051 | False |
| Event Driven | 0.006344 | 0.0084 | False |
| Fixed Income Arbitrage | 0.004365 | 0.0055 | False |
| Global Macro | 0.005403 | 0.0038 | True |
| Long/Short Equity | 0.006331 | 0.0079 | False |
| Merger Arbitrage | 0.005356 | 0.0060 | False |
| Relative Value | 0.005792 | 0.0067 | False |
| Short Selling | -0.001701 | -0.0053 | True |
| Funds Of Funds | 0.004262 | 0.0052 | False |

Now, let's develop the code to compute the skewness of a series of numbers.

Recall that the skewness is given by:

$$S(R) = \frac{E[(R - E(R))^3]}{\sigma_R^3}$$

```
1 def skewness(r):
2     """
3     Alternative to scipy.stats.skew()
4     Computes the skewness of the supplied Series or DataFrame
5     Returns a float or a Series
6     """
7     demeaned_r = r - r.mean()
8     # use the population standard deviation, so set dof=0
9     sigma_r = r.std(ddof=0)
10    exp = (demeaned_r**3).mean()
11    return exp/sigma_r**3
12
```

```
1 skewness(hfi).sort_values()
```

```
↳ Fixed Income Arbitrage    -3.940320
   Convertible Arbitrage    -2.639592
   Equity Market Neutral    -2.124435
   Relative Value           -1.815470
   Event Driven             -1.409154
   Merger Arbitrage         -1.320083
   Distressed Securities    -1.300842
   Emerging Markets        -1.167067
   Long/Short Equity        -0.390227
   Funds Of Funds          -0.361783
   CTA Global               0.173699
   Short Selling            0.767975
   Global Macro             0.982922
dtype: float64
```

Just to see if we get the same answer, let's use the skewness function that is built into `scipy.stats`

```
1 import scipy.stats
2 scipy.stats.skew(hfi)
```

```
↳ array([-2.63959223,  0.17369864, -1.30084204, -1.16706749, -2.12443538,
        -1.40915356, -3.94032029,  0.98292188, -0.39022677, -1.32008333,
        -1.81546975,  0.76797484, -0.36178308])
```

So, let's add that to our `edhec_risk_kit.py`.

Finally, let's look at the skewness that you would expect from a truly random sequence of returns. I'll create a generator from numpy and generate the same number of returns as we have for the hedge fund data.

```
1 hfi.shape
```

```
↳ (263, 13)
```

```
1 import numpy as np
2 normal_rets = np.random.normal(0, 0.15, (263, 1))
```

```
1 normal_rets.mean(), normal_rets.std()
```

```
↳ (9.807196380857325e-05, 0.15552295012928952)
```

```
1 erk.skewness(normal_rets)
```

```
↳ -0.030687276528881435
```

▼ Kurtosis

Intuitively, the kurtosis measures the "fatness" of the tails of the distribution. The normal distribution has a kurtosis of 3. If the kurtosis of your returns is less than 3 then it tends to have thinner tails, and if the kurtosis is greater than 3 then it tends to have fatter tails.

Kurtosis is given by:

$$K(R) = \frac{E[(R - E(R))^4]}{\sigma_R^4}$$

This is very similar to the skewness, so we can just copy and paste it and then edit it to compute the kurtosis (was the case for skewness).

```
1 erk.kurtosis(hfi)
```

```
↳ Convertible Arbitrage      23.280834
   CTA Global                 2.952960
   Distressed Securities      7.889983
   Emerging Markets          9.250788
   Equity Market Neutral     17.218555
   Event Driven               8.035828
   Fixed Income Arbitrage    29.842199
   Global Macro               5.741679
   Long/Short Equity         4.523893
   Merger Arbitrage          8.738950
   Relative Value            12.121208
   Short Selling              6.117772
   Funds Of Funds            7.070153
dtype: float64
```

Let's compare it with `scipy.stats` ...

```
1 scipy.stats.kurtosis(hfi)
```

```
↳ array([20.28083446, -0.04703963,  4.88998336,  6.25078841, 14.21855526,
         5.03582817, 26.84219928,  2.74167945,  1.52389258,  5.73894979,
         9.12120787,  3.11777175,  4.07015278])
```

Note that these numbers are all lower by 3 from the number we have computed. That's because, as we saw, the kurtosis of a normally distributed series of numbers is 3, and `scipy.stats` is returning the *Excess Kurtosis* (the kurtosis minus 3). Applying it on the random normal numbers we generated:

```
1 scipy.stats.kurtosis(normal_rets)
```

```
↳ array([-0.27076178])
```

```
1 erk.kurtosis(normal_rets)
```

```
↳ 2.729238217614401
```

▼ Running the Jarque-Bera Test for Normality

The `scipy.stats` module contains a function that runs the *Jarque-Bera* test on a sequence of normally generated returns:

```
1 scipy.stats.jarque_bera(normal_rets)
```

```
↳ (0.8446549484268455, 0.6555193385610887)
```

The first number is the test statistic and the second number is the one we want. It represents the p-value. If you want to run the test at a 1% level of significance, you want this number to be greater than 0.01 to accept the hypothesis of normality, and if that number is less than 0.01 then you must reject the hypothesis of normality.

In this case, since we got a number higher than 0.01 we can accept the hypothesis that the numbers are normally distributed for our different hedge fund indices.

```
1 scipy.stats.jarque_bera(hfi)
```

```
↳ (25656.585999171326, 0.0)
```

Why didn't we get the results for the individual indices? Because the implementation of the test isn't designed to treat each column as a separate set of returns. We can write our own wrapper for it to fix this, and adding this code to our python file:

```
import scipy.stats

def is_normal(r, level=0.01):
    """
    Applies the Jarque-Bera test to determine if a Series is normal or not
    Test is applied at the 1% level by default
    Returns True if the hypothesis of normality is accepted, False otherwise
    """
    statistic, p_value = scipy.stats.jarque_bera(r)
    return p_value > level
```

```
1 erk.is_normal(normal_rets)
```

```
☞ True
```

There are a few different ways to handle the problem. The first is to use the `.aggregate` method on `hfi` as an argument and applies that function to each column:

```
1 hfi.aggregate(erk.is_normal)
```

```
☞ Convertible Arbitrage      False
   CTA Global                 True
   Distressed Securities      False
   Emerging Markets          False
   Equity Market Neutral      False
   Event Driven               False
   Fixed Income Arbitrage     False
   Global Macro               False
   Long/Short Equity          False
   Merger Arbitrage           False
   Relative Value             False
   Short Selling              False
   Funds Of Funds            False
   dtype: bool
```


However, we can fix this in our wrapper so that we have a uniform interface to test normality:

```
import scipy.stats

def is_normal(r, level=0.01):
    """
    Applies the Jarque-Bera test to determine if a Series is normal or not
    Test is applied at the 1% level by default
    Returns True if the hypothesis of normality is accepted, False otherwise
    """
    if isinstance(r, pd.DataFrame):
        return r.aggregate(is_normal)
    else:
        statistic, p_value = scipy.stats.jarque_bera(r)
        return p_value > level
```

```
1 import pandas as pd
2 isinstance(hfi, pd.DataFrame)
```

```
☞ True
```

```
1 erk.is_normal(normal_rets)
```

```
☞ True
```

▼ Testing CRSP SmallCap and Large Cap returns for Normality

Let's see whether any of the returns we've been studying so far pass the normality hypothesis.

```
1 ffme = erk.get_ffme_returns(DATAPATH)
2 erk.skewness(ffme)
```

```
☞ SmallCap    4.410739
   LargeCap    0.233445
   dtype: float64
```

```
1 erk.kurtosis(ffme)
```

```
☞ SmallCap    46.845008
   LargeCap    10.694654
   dtype: float64
```

```
1 erk.is_normal(ffme)
```

```
☐→ SmallCap    False  
    LargeCap    False  
    dtype: bool
```

```
1
```